

kaggle

Search

Competitions

Datasets

Notebooks

Discussion

Courses

...



Code

This Notebook has been released under the [Apache 2.0](#) open source license.

Download Code

```

1  ##### Introduction #####
2  # In the past, the best submissions for this data mining challenge
3  # (the challenge not the data set) have been simple logistic
4  # regression models that use variables that make intuitive sense. It is absolutely paramount for a data
5  # scientist to understand the domain the data is from. There are examples where 10 Phd guys build
6  # some ridiculous NN and then there is this guy who regresses three variables out of the 100 or so
7  # and comes up with a better model.
8  # Having said that, I am lazy and intrigued at the same time. I don't want to dig into the data
9  # too much, but just want to build a classifier for the hell of it. Classic mistake.
10 # There are two reasons why in this particular case this is a bad idea. First, the data is not
11 # very complex. Second, it is a very intuitive case. If these conditions are met, it is likely better
12 # to build a simple classifier, instead of going ham on the methods that are available, like
13 # building a random forest. This approach generally means you have no idea what you're dealing
14 # with and/or are just enjoying the academic challenge.
15 # This should be the lesson of this competition, KISS - Keep It Simple Stupid. Too bad it won't stick.
16 # So let's go down the rabbit hole.
17 require(Matrix)
18 require(data.table)
19
20 # read data
21 d <- read.csv('../input/carInsurance_train.csv')
22 # classify <- read.csv('~/Desktop/ML/DSS_DMC_2017/Working Material/test_set.csv')
23
24 ##### data exploration #####
25
26 ## should probably do outliers and data exploration
27 boxplot(d) # wow, balance. Could think about normalizing it and kicking the guy up top out. Let's check
28 # if balance correlates with the Car Insurance to be classified
29
30 cor.test(d$Balance, d$CarInsurance)
31 # pretty significant. What if we remove some "outliers"?
32 cor.test(d[d$Balance <= 20000, "Balance"], d$CarInsurance[d$Balance <= 20000])
33 # stops being significant. I am definitely overlooking some mathematical truth that would be apparent if I took the
34
35 ## let's look at the others without balance throwing everything in disarray
36 boxplot(d[, -c(1,7)]) # obviously we don't need ID either
37 # days passed still messed up
38
39 cor.test(d$DaysPassed, d$CarInsurance)
40 # pretty relevant
41
42 ##### We'll leave it at that. It should give you an idea. Caret is a good resource as well http://topepo.github.io/c
43
44 ##### Cleaning #####
45
46 ## Cleaning has two parts. Getting rid of the usual, like handling NAs and dealing with nonsensical stuff (e.g. ID
47 ## My assumptions are that the day and time of day for the calls influence conversion rates.
48 # ***** should probably look at the columns for LastCallDay and LastCallMonth. They don't make sense in their cur
49 # Actually, finding out which weekday calls were made to what time would be very helpful. Because cold calls tend t
50 # We start with the 2nd part first, because NA imputation also somewhat depends on the steps we will take here. You
51
52 ## Get the weekdays in a workable format
53 # What are we doing here? We believe there is a chance that calls made at a certain time on certain days, e.g. Thur
54 # However, we don't know which year this data is from. Which does not matter. We just need to have a uniform distri

```

```

56
57 clean_d <- d
58
59 max(clean_d[clean_d$LastContactMonth == 'feb', 'LastContactDay']) # check if the year the calls were made isn't a leap year
60 clean_d$dateCall <- as.Date(paste(clean_d$LastContactDay, clean_d$LastContactMonth, "2015", sep = '/'), "%d/%b/%Y")
61 clean_d$Weekday <- factor(weekdays(clean_d$dateCall))
62
63 # Next, we want to know what time people were called during the day.
64 # Let's see when the calls were made and what the working hours are.
65 plot(table(clean_d$CallStart)) # not very informative let's take the minutes and seconds off
66 plot(table(call_hr <- gsub("(:\\d{2})", "", clean_d$CallStart))) # ok... they are pretty diligent in calling people
67
68 # We could take the times as they are given. However, that would be too much noise, in my opinion. Therefore, I opt for a more holistic
69 clean_d$CallDayTime <- as.numeric(gsub("(:\\d{2})", "", clean_d$CallStart))
70 require(car)
71 clean_d$CallDayTime <- factor(recode(clean_d$CallDayTime, "c('9', '10', '11')='morning'; c('12', '13', '14')='midday'; c('15', '16', '17')='evening'"))
72
73 ##
74 ##
75 # Convert last phone call time to minutes
76 require(chron)
77 clean_d$call_dur_min <- 60 * 24 * as.numeric(times(clean_d$CallEnd)-times(clean_d$CallStart))
78
79 ## find NAs
80 # I didn't come up with that. Google + Stackoverflow = Bliss
81 na_count <- sapply(clean_d, function(y) sum(length(which(is.na(y)))))
82 na_count <- data.frame(na_count)
83
84 ## The issue is that we have a lot of NAs. Question is, what to do with them? Kick the missing values? Include them?
85 # But most classifiers can't handle missing values. So we should replace them. For a more holistic
86 # overview of the how's and why's read the article above. We are going to go with the cool solution.
87 # We replace NAs with factors falling out of a k-Nearest-Neighbor
88
89 ## Before we can do that, we need to take away all the columns that we won't consider in our analysis. Reason being
90 summary(clean_d)
91 # ID is the first to go, redundant a.f.. Outcome probably too much noise, has to go. CallStart and CallEnd, LastContactDay, LastContactMonth
92 # Kick everything we don't need. This is a little complicated in R if you want to do it by column name. We do it another way
93
94 sub_clean_d <- subset(clean_d, select = -c(Id, LastContactDay, LastContactMonth, Outcome, CallStart, CallEnd, DateCall))
95
96 # We can do the NA replacement now
97 ## Impute NAs https://www.r-bloggers.com/missing-value-treatment/
98 # this might leave us in Hell's Kitchen. Well... actually it's Freedman's Kitchen https://www.r-bloggers.com/freedmans-kitchen/
99
100 require(DMwR)
101 set.seed(42)
102
103 clean_d_imputeknn <- knnImputation(sub_clean_d) # perform knn imputation.
104 # check if it missed any NAs
105 anyNA(clean_d_imputeknn)
106
107 #####
108 ##### Building the actual model #####
109 ## We start with a random forest.
110 ## Using the caret package we cross validate before we train the model
111 # Cross validation is pretty much the most important thing, because it reduces overfitting. It means we partition the data into
112 # Basically with the caret package, we define a cross validation object which we pass to the actual function to create the model
113
114 ## Making a nice data frame for the modelling operations.
115 model_d <- clean_d_imputeknn
116
117
118
119 # Split data into test and train: so later we can validate our model
120 require(caret)
121 set.seed(42)
122
123 # We will use this index to define the rows, which will go into train and test data sets respectively
124 train_index <- createDataPartition(model_d$CarInsurance, p = 0.75, list = FALSE, times = 1)
125 training <- model_d[train_index, ]

```



Simple Random Forest on Insurance Call Forecast

R script using data from [Car Insurance Cold Calls](#) · 985 views · 2y ago

0

Copy and Edit

6

Ver
7

```

130
131
132 ##### Random Forest #####
133 # Training the actual model. We have to pass CarInsurance, i.e. the variable to be classified, as a factor, not as
134 set.seed(42)
135 model_rf = train(factor(CarInsurance)~., data=training, trControl=train_control, method="rf")
136
137 # We make a frame and fill it with the predicted values. This allows us to the test the quality of the model in the
138 prediction_rf = predict(model_rf, subset(testing, select=-c(CarInsurance)))
139
140 #Compute the accuracy of predictions with a confusion matrix
141 confusionMatrix(prediction_rf, testing$CarInsurance)
142
143 ##### Logistic Regression #####
144 ## simple first, then bagging and boosting
145 set.seed(42)
146 model_logreg <- glm(factor(CarInsurance) ~., family=binomial(link='logit'), data=training)
147
148 prediction_logreg = predict(model_logreg, subset(testing, select=-c(CarInsurance)), type='response') # by choosing
149
150 table(testing$CarInsurance, prediction_logreg > 0.5) # LogReg gives the results as probabilities, so we can't use t
151
152 ## so Random Forest results are actually better. However, the RF takes about a minute or so to calculate the model.
153
154 ##### LogitBoost #####
155 set.seed(42)
156 model_logitboost <- train(factor(CarInsurance)~., data=training, trainControl=train_control, method="LogitBoost", n
157
158 prediction_logitboost = predict(model_logitboost, subset(testing, select=-c(CarInsurance)))
159 confusionMatrix(prediction_logitboost, testing$CarInsurance)
160
161 ##### XGB Trees #####
162 set.seed(42)
163 model_xgbtrees <- train(factor(CarInsurance)~., data=training, method='xgbTree', trainControl=train_control, metric
164
165 prediction_xgbtrees = predict(model_xgbtrees, subset(testing, select= -c(CarInsurance)))
166 confusionMatrix(prediction_xgbtrees, testing$CarInsurance) # worst result yet :(

```

Did you find this Notebook useful?
Show your appreciation with an upvote

▲
0

Run Info

Succeeded	False	Run Time	1200.5 seconds
Exit Code	137	Queue Time	0 seconds
Docker Image Name	kaggle/rstats (Dockerfile)	Output Size	0
Timeout Exceeded	True	Used All Space	False
Failure Message	The kernel was killed for running longer than 1200 seconds.		

Log

Download Log

Time Line # Log Message



Code



Log



Data



Output



Comments

```
3.1s      4      Pearson's product-moment correlation

data: d$Balance and d$CarInsurance
3.1s      5      t = 2.6302, df = 3998, p-value = 0.008567
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.01058319 0.07245913
sample estimates:
cor
0.04156101

      Pearson's product-moment correlation

data: d[d$Balance <= 20000, "Balance"] and d$CarInsurance[d$Balance <= 20000]
3.1s      7      t = 4.2517, df = 3974, p-value = 2.17e-05
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.03628313 0.09817159
sample estimates:
cor
0.06729209

3.2s      8      Pearson's product-moment correlation

data: d$DaysPassed and d$CarInsurance
3.2s      9      t = 8.8714, df = 3998, p-value < 2.2e-16
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.1084183 0.1692057
sample estimates:
cor
0.1389429

[1] 27
3.3s     10      Loading required package: car
4.1s     11      Loading required package: chron
4.5s     12      Id      Age      Job      Marital
Min. : 1      Min. :18.00      management :893      divorced: 483
1st Qu.:1001      1st Qu.:32.00      blue-collar:759      married :2304
Median :2000      Median :39.00      technician :660      single  :1213
Mean :2000      Mean :41.21      admin. :459
3rd Qu.:3000      3rd Qu.:49.00      services :330
4.5s     13      Max. :4000      Max. :95.00      (Other) :880
NA's :19

      Education      Default      Balance      HHInsurance
primary : 561      Min. :0.0000      Min. : -3058.0      Min. :0.0000
secondary:1988      1st Qu.:0.0000      1st Qu.: 111.0      1st Qu.:0.0000
tertiary :1282      Median :0.0000      Median : 551.5      Median :0.0000
NA's : 169      Mean :0.0145      Mean : 1532.9      Mean :0.4928
3rd Qu.:0.0000      3rd Qu.:1619.0      3rd Qu.:1.0000
Max. :1.0000      Max. :98417.0      Max. :1.0000

      CarLoan      Communication      LastContactDay      LastContactMonth
Min. :0.000      cellular :2831      Min. : 1.00      may :1049
1st Qu.:0.000      telephone: 267      1st Qu.: 8.00      jul : 573
Median :0.000      NA's : 902      Median :16.00      aug : 536
Mean :0.133      Mean :15.72      Mean :15.72      jun : 454
3rd Qu.:0.000      3rd Qu.:22.00      3rd Qu.:22.00      nov : 347
Max. :1.000      Max. :31.00      Max. :31.00      apr : 306
(Other): 735

      NoOfContacts      DaysPassed      PrevAttempts      Outcome
Min. : 1.000      Min. : -1.00      Min. : 0.0000      failure: 437
1st Qu.: 1.000      1st Qu.: -1.00      1st Qu.: 0.0000      other : 195
Median : 2.000      Median : -1.00      Median : 0.0000      success: 326
Mean : 2.607      Mean : 48.71      Mean : 0.7175      NA's :3042
3rd Qu.: 3.000      3rd Qu.: -1.00      3rd Qu.: 0.0000
Max. :43.000      Max. :854.00      Max. :58.0000

      CallStart      CallEnd      CarInsurance      DateCall
10:42:44: 3      10:22:30: 3      Min. :0.000      Min. :2015-01-08
11:48:25: 3      10:52:24: 3      1st Qu.:0.000      1st Qu.:2015-05-08
13:54:34: 3      11:27:46: 3      Median :0.000      Median :2015-06-05
15:27:56: 3      09:04:02: 2      Mean :0.401      Mean :2015-06-21
15:48:27: 3      09:06:42: 2      3rd Qu.:1.000      3rd Qu.:2015-08-11
17:02:39: 3      09:12:47: 2      Max. :1.000      Max. :2015-12-30
(Other) :3982      (Other) :3985

      Weekday      CallDayTime      call_dur_min
Friday :725      afternoon:1351      Min. : 0.08333
Monday :454      midday :1365      1st Qu.: 2.10000
Saturday:380      morning :1284      Median : 3.86667
Sunday : 92      Mean : 5.84740
Thursday :849      3rd Qu.: 7.66667
Tuesday :724      Max. :54.21667
Wednesday:776
```

```

4.5s      14 Loading required package: DMwR
4.6s      15 Loading required package: methods
4.7s      16 Loading required package: lattice
4.7s      17 Loading required package: grid
45.8s     18 [1] FALSE
45.8s     19 Loading required package: caret
45.8s     20 Loading required package: ggplot2
46.2s     21 Loading required package: randomForest
46.2s     22 randomForest 4.6-12
Type rfNews() to see new features/changes/bug fixes.
46.3s     23
Attaching package: 'randomForest'

The following object is masked from 'package:ggplot2':

    margin

197.6s    24 Confusion Matrix and Statistics

              Reference
Prediction   0     1
0      483    79
1      106   332

197.6s    25

              Accuracy : 0.815
              95% CI   : (0.7895, 0.8386)
              No Information Rate : 0.589
              P-Value [Acc > NIR] : < 2e-16

              Kappa   : 0.6216
              Mcnemar's Test P-Value : 0.05593

              Sensitivity : 0.8200
              Specificity : 0.8078
              Pos Pred Value : 0.8594
              Neg Pred Value : 0.7580
              Prevalence : 0.5890
              Detection Rate : 0.4830
              Detection Prevalence : 0.5620
              Balanced Accuracy : 0.8139

              'Positive' Class : 0

197.7s    26

              FALSE TRUE
0      499    90
1      117   294

198.3s    27 Loading required package: caTools
211.5s    28 Confusion Matrix and Statistics

              Reference
Prediction   0     1
0      506   147
1       83   264

211.5s    29

              Accuracy : 0.77
              95% CI   : (0.7426, 0.7958)
              No Information Rate : 0.589
              P-Value [Acc > NIR] : < 2.2e-16

              Kappa   : 0.5135
              Mcnemar's Test P-Value : 3.266e-05

              Sensitivity : 0.8591
              Specificity : 0.6423
              Pos Pred Value : 0.7749
              Neg Pred Value : 0.7608
              Prevalence : 0.5890
              Detection Rate : 0.5060
              Detection Prevalence : 0.6530
              Balanced Accuracy : 0.7507

              'Positive' Class : 0

211.6s    30 Loading required package: xgboost
211.6s    31 Loading required package: plyr
211.6s    32
Attaching package: 'plyr'

The following object is masked from 'package:DMwR':

    join

211.6s    33
211.6s    35 Failed. Exited with code 137.

```

Data

Data Sources

▼ Car Insurance Cold Calls

carInsurance_test.csv	19 columns
carInsurance_train.csv	19 columns
DSS_DMC_Description.pdf	



Car Insurance Cold Calls

We help the guys and girls at the front to get out of Cold Call Hell

Last Updated: 2 years ago (Version 1)

About this Dataset

Introduction

Here you find a very simple, beginner-friendly data set. No sparse matrices, no fancy tools needed to understand what's going on. Just a couple of rows and columns. Super simple stuff. As explained below, this data set is used for a competition. As it turns out, this competition tends to reveal a common truth in data science: KISS - Keep It Simple Stupid

What is so special about this data set is, given it's simplicity, it pays off to use "simple" classifiers as well. This year's competition was won by a C5.0 . Can you do better?

Description

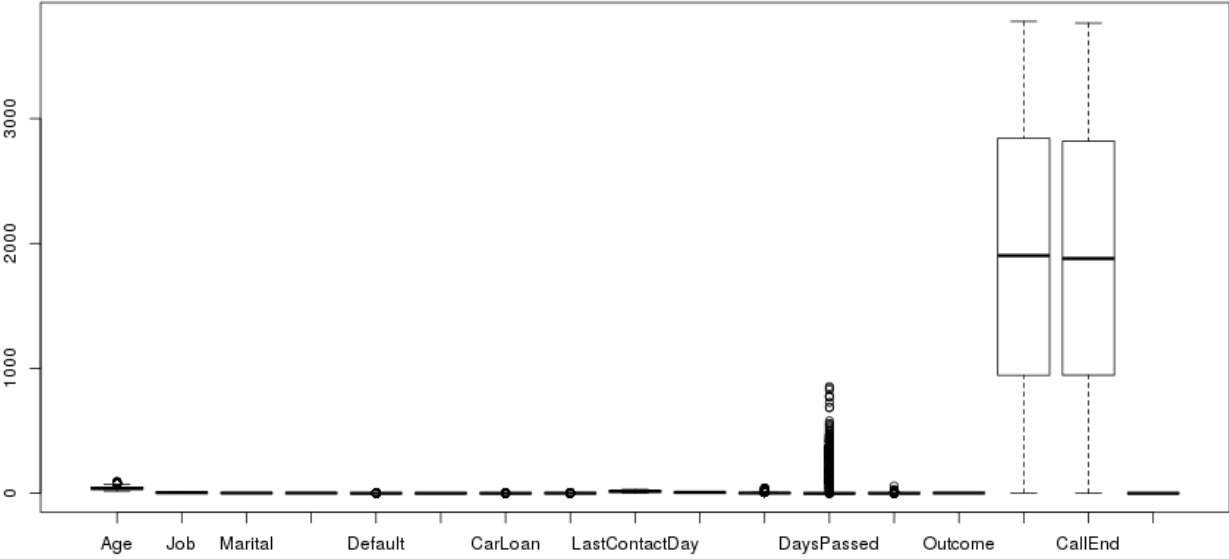

We are looking at cold call results. Turns out, same salespeople called existing insurance customers up and tried to sell car insurance. What you have are details about the called customers. Their age, job, marital status, whether the have home insurance, a car loan, etc. As I said, super simple.

What I would love to see is some of you applying some crazy XGBoost classifiers, which we can square off against some logistic regressions. It would be curious to see what comes out on top. Thank you for your time, I hope you enjoy using the data set.

Acknowledgements

Thanks goes to the Decision Science and Systems Chair of Technical University of Munich (TUM) for getting the data set

Output Visualizations





Comments (0)



Click here to comment...